

HARSHIBAR

555.555.5555

hello@email.com

linkedin.com/in/harshibar

github.com/harshibar

U.S. Citizen

SUMMARY

Data Scientist with experience building end-to-end machine learning systems, from data pipelines and feature engineering to model deployment and monitoring. Strong background in statistical modeling, NLP/LLMs, and turning ambiguous business questions into measurable outcomes.

EXPERIENCE

Google Verily

Aug. 2018 – Sept. 2019

Data Scientist

San Francisco, CA

- Built ML models to flag early-stage cancer signals across 50k blood samples, improving detection recall by 18% over the prior rules-based baseline
- Designed a feature store and reproducible training pipeline with 20 cross-functional stakeholders, cutting model iteration time from weeks to days and saving \$1M annually
- Ran controlled experiments (A/B tests, causal analysis) on a lab workflow tool, driving a 40% efficiency gain validated across 10 core users

Amazon

May 2017 – Aug. 2017

Data Science Intern

Seattle, WA

- Joined the Search Customer Experience Team; received a return offer for a full-time role
- Shipped a ranking-model improvement to 2M+ users that lifted click-through on movie-series queries by 12%
- Built a back-end data service in Java/Python to log and analyze query behavior, enabling future model-training datasets

YouTube

Aug. 2019 – Present

Creator & Independent Data Analyst (@harshibar)

San Francisco, CA

- Grew a tech channel to 60k subscribers in 1.5 years; analyzed viewership data to guide content strategy
- Ran A/B tests on titles and thumbnails, increasing impressions by 2.5M in 3 months using data-driven iteration

PROJECTS

RAG-Powered Document Assistant | *Python, LangChain, OpenAI, FAISS, Streamlit*

- Built a retrieval-augmented-generation chatbot that answers questions over private PDFs using vector embeddings and semantic search, reducing manual document lookup time by 70%
- Implemented chunking, embedding, and a re-ranking layer to cut hallucinations; deployed an interactive Streamlit front end

End-to-End ML Pipeline with MLOps | *scikit-learn, MLflow, Docker, FastAPI, AWS*

- Developed a churn-prediction model (XGBoost, AUC 0.89) with automated feature engineering and hyperparameter tuning
- Containerized the model behind a FastAPI endpoint, tracked experiments in MLflow, and set up drift monitoring for continuous retraining

Time-Series Demand Forecasting Dashboard | *Python, Prophet, Pandas, Plotly*

- Forecasted retail demand across 500+ SKUs using Prophet and LSTM baselines, lowering forecast error (MAPE) by 22% versus the naive seasonal model
- Shipped an interactive Plotly dashboard for stakeholders to explore predictions, confidence intervals, and scenario planning

EDUCATION

Wellesley College

Aug. 2014 – May 2018

Bachelor of Arts in Computer Science

Wellesley, MA

- Coursework: Machine Learning, Data Structures, Algorithms, Databases, Statistics
- Research: MIT Graybiel Lab (published author); MIT Media Lab (analyzed urban microbe spread)

SKILLS

Languages: Python, SQL (PostgreSQL, MySQL), R, JavaScript

ML / Data: scikit-learn, PyTorch, TensorFlow, XGBoost, Pandas, NumPy, LangChain

MLOps / Tools: MLflow, Docker, FastAPI, AWS, Git, Airflow

Visualization: Plotly, Tableau, Matplotlib, Seaborn, Google Analytics